

Cassandra & Akka

IndyScala, August 5, 2013

Ross A. Baker, Sr. Cloud Engineer



What is Cassandra?

- A persistent Map[A, SortedMap[B, C]]*
- distributed
- fault tolerant
- tunably consistent
- FAST!

** Oversimplification. Bear with me.*

Let's build a system that...

1. Ingests the Twitter Sample API*
2. Detects the language of the tweets
3. Create a searchable index, by language
4. Records info about users:
 - a. Their interests (#hashtags)
 - b. People who talk to them (@mentions)
 - c. The languages they tweet in

** I'm not worthy of the firehose*

Our schema

tweets	
tweet_id (PK)	bigint
text	text
username	text
lang	map<text, text>

user_profiles	
username	text
languages	set<text>
hashtags	set<text>
mentions	set<text>

tweet_terms	
term (PK)	text
lang (PK)	text
tweet_id (PK)	bigint
username	text
text	text

counters	
name	count

How do you define it?

CQL3 DDL looks a lot like SQL DDL:

```
create table tweets (  
    tweet_id bigint primary key,  
    text text,  
    username text,  
    lang map<text, text>  
);
```

How do you query it?

CQL3 DML looks a lot like SQL DML:

```
INSERT INTO tweets (tweet_id, text, username)  
VALUES (?, ?, ?);
```

```
SELECT * FROM tweets WHERE tweet_id = ?;
```

But it's not relational

- No check constraints on columns
- No foreign keys to other tables
- Not normalized.
- Not even in 1NF!!!

So it doesn't do...

- joins
- subqueries
- like queries
- most forms of aggregation
- filtering by non-PK columns*
- range queries on PK**
- ordering by column value***

* without a secondary index

** you can, with an inefficient config

*** unless it's the second column of the primary key

That sucks!

- Yeah, at first, if you're a relational guy, it does suck.
- Think in terms of the data structure.
 - Review: Map[A, SortedMap[B, C]]
- Think in terms of how query that structure.*
- Don't fear denormalization.
- You don't have to give up your other DBMS.

* If you take away one thing from this whole evening, take that.

Seems limiting...

But it's also freeing.

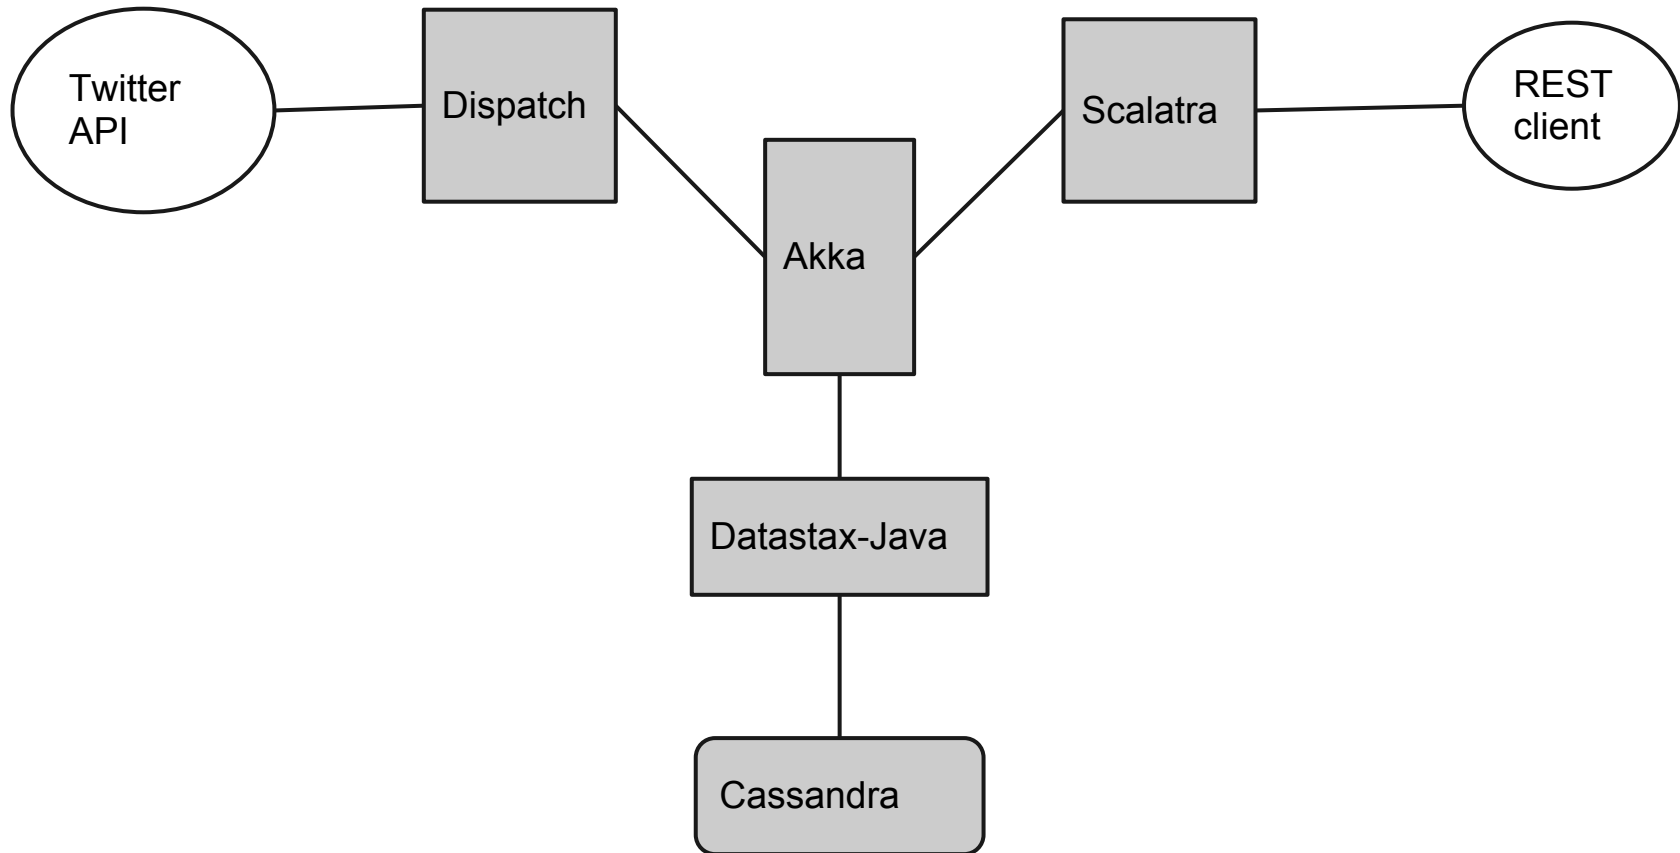
- Sharding is intrinsic.
- Replication is intrinsic.
- Add nodes dynamically. *
- Available if some nodes fail. **
- Available if Godzilla rampages through one of your data centers. ***

* Not a cheap operation, but no downtime.

** How many? Depends on your *replication factor*.

*** If you set up replication across centers, and if Godzilla doesn't get them all.

Architecture



Walk through the code

Open your browsers and follow along as we discuss the important classes.

<https://github.com/indyscala/cassandra-demo/tree/master/src/main/scala/indyscala/aug2013>

source.TwitterSample

- Uses Dispatch to connect to the Twitter Sample API.
- Parses each line of the streaming response with json4s.
- Executes a callback function. In our case, we send it to the Twitter Service.

service.TweetService

- Our central nervous system.
- Supervises our subservices:
 - tweet parser actor
 - data store actor
 - language detector actor
 - indexer actor
- Knows how to route:
 - incoming JSON messages from the TwitterSample
 - parsed tweets from the
 - incoming queries from Scalatra
 - results from the language detection actor
 - results from the index actor

service.JsonParser

- Knows how to construct a tweet from a JSON value.
- Actually, the smarts are built into extractors in the model package.
- Important point: feeds recognized Tweets back into the TweetService.

service.LangDetector

- Detects languages from two sources
 - Twitter gives us its best guess
 - Analyzes with Cybozu Labs' LangDetect, an open source project.
- Does the detections in parallel, and reports back to the TweetService to store the results and trigger indexing.
- Adding an HTTP services for a third opinion via Dispatch would be easy.
 - 10% of my speaker fee of \$0 bought us 0 such services.

service.Indexer

- Triggered when:
 - a tweet is parsed (for any-language searches)
 - a language is detected
- Yes, yes, caching parsed tweets may have been more efficient. But this approach was simple.
- Yes, yes, the index table is repetitious. But it's efficient in Cassandra. Forget your relational instincts!

service.DataStore

- Finally... this actor is why we're all gathered here tonight.
- All Cassandra interaction happens in this actor.
 - That's not necessary. I have lots of actors that talk to Cassandra at work. This is a simple app.
- It's all non-blocking. Great fit for Akka.
- DataStax's Java driver is nice, but you deal with some types you left behind when you left Java behind...

util.DataStaxImplicits

- Implicits to the rescue!
- Adapt Guava's Listenable Future to scala.concurrent.Future, makes working with Cassandra feel like a first-class citizen in Akka.
- Map ResultSets over an arbitrary function, to feel more like a Scala collection.
 - Map straight to json with json4s, because that's what we usually do.

util.DataStaxExtension

- Provides a way for multiple actors to share a Cassandra connection pool.
 - Only one actor type talks to Cassandra in this demo, but we have more than one instance of it.
- It's convenient like a global, but it's scoped to the ActorSystem.
 - We only have one ActorSystem, but you'll really appreciate this when you start writing tests.

servlet.TwitterServlet

- A simple Scalatra service to define a few REST endpoints and query the service.
- Sends query messages to the TwitterService, and renders the returned JSON.
- Yes, Virginia, there is asynchronous support in the Servlet API.
 - Are you the same person you were in 1997? Neither are Java Servlets.

What still hurts

- Typing
 - Known invalid CQL3 types can be bound. In native Scala, we'd have a type class here.
 - Collections are Java collections, and because of the Any type, we have to convert them before binding.
- Positional parameters. Eek!
- Anorm has done a nice job applying a Scala sheen to raw JDBC. Something similar for talking to Cassandra would be a big win with little added complexity.

Immaturity

- Don't panic. I've not lost any data yet.
 - But some things aren't implemented yet. Like prepared statements with a dynamic limit. Sometimes we're stuck with *gasp* String concat.
- Some of the immaturity is in me.
 - A dozen years of RDBMS wisdom was hard earned, and picking Cassandra means starting over.
- Some of the immaturity is in us all.
 - Nobody has a dozen years of Cassandra wisdom, because it's not a dozen years old. We're all learning as we go.

In conclusion

- I work with Big Data. Cassandra is doing things for us that PostgreSQL simply can't.
- But sometimes I need ad hoc queries on smaller data sets, and PostgreSQL is still really great at that.
- It's okay to use them both. *NoSQL means Not Only SQL.*
 - And with Akka and future comprehensions, you've got a great toolset to join them in your application.

Does this look like fun?

- I work with Cassandra, Akka, Json4s, Scalatra, and Dispatch every day at CrowdStrike.
- We're a distributed team.
- And we're hiring!



CROWDSTRIKE

Thanks and good night

- Thanks to E-gineering for hosting.
- Thanks to Brad Fritz (@bfritz) of Fewer Hassles for organizing
- Once again, I'm @rossabaker
 - Also rossabaker on Freenode IRC. I'd love to chat.
- Materials: <http://github.com/indyscala/cassandra-demo>